

Recurrent networks

Dataset: Chicago transit ridership

How to use these notes

The primary text for this lecture is Géron, Chapter 13. These notes are not a replacement for it and do not repeat it: they are the lecture’s own argument written out at length — the derivation done slowly, the numbers with the runs that produced them, and the reasoning between one slide and the next that a deck can only gesture at. Read the chapter for the method; read these for why the lecture makes the choices it does. Every figure quoted here is printed by this lecture’s notebook, and the course checks that mechanically rather than trusting it.

Contents

1	What a linear model on lags cannot do	2
1.1	Unrolling, and the problem it creates	2
2	Deeper, wider, and further ahead	4
2.1	Loss, scale and what to watch	4
2.2	Forecasting k steps ahead	4
2.3	Improvement 1 — deeper, which this run cannot rank	5
2.4	Improvement 2 — more series, which worked	5
2.5	Improvement 3 — further than one day	6
3	LSTM and GRU, gate by gate	6
4	The whole application	7
4.1	Red-team it: what happens when the world changes	8
5	The notebook	9
6	Exercises	10

1 What a linear model on lags cannot do

Lecture 15 established the ground rules and a number to beat: evaluation is forward in time, never a random split; the seasonal-naive forecast is genuinely hard to beat and the autocorrelation says why; and a linear model on lagged features beats it, modestly.

That model has a *fixed window*. Anything outside the last 56 days is invisible to it, and every lag gets one weight regardless of context.

A recurrent network keeps a *state* instead, and that state is not bounded by a window:

$$\mathbf{h}_t = \phi\left(\mathbf{W}_x^T \mathbf{x}_t + \mathbf{W}_h^T \mathbf{h}_{t-1} + \mathbf{b}\right).$$

This is Lecture 9's $\phi(\mathbf{XW} + \mathbf{b})$ with one extra term, and that term is the whole idea. $\mathbf{W}_x$ and $\mathbf{W}_h$ do not depend on t : one set of weights applied at every step, so reach is decoupled from parameter count — the same argument as weight sharing in Lecture 12, along time instead of across an image.

The state is a summary, and that is a budget

$\mathbf{h}_t$ is everything the network has decided to remember about $x_1, \dots, x_t$, at *fixed width* whatever the sequence length — so it is lossy by construction. What it keeps is learned, not designed: nothing tells the cell that “the same weekday last week” matters.

Two sequences with the same state are indistinguishable to everything downstream, which is the definition of a sufficient statistic, and the network is trying to learn one. It is also the honest limit: if the task needs more bits than the state has, no amount of training supplies them.

1.1 Unrolling, and the problem it creates

Draw the same cell once per time step and you get a network 56 layers deep with one set of weights, reused. Then differentiate it exactly as Lecture 10 taught — **backpropagation through time** needs nothing new.

Failure condition — one matrix raised to the length of the sequence

Unrolling 56 steps gives a 56-layer network, and Lecture 11 said what happens to a gradient descending 56 layers: it is multiplied by the same factor 56 times. Here the factor is *the same matrix* every step, because the weights are shared — so the gradient is governed by a power of one matrix, and powers of a matrix either vanish or explode unless the spectral radius sits almost exactly at one.

Initialisation cannot fix this the way it did in Lecture 11, where the layers were independent. Something has to change in the cell itself.

Shape	Input	Output	Example
sequence to vector	a sequence	one value	today's model; classify a review
vector to sequence	one value	a sequence	a caption from an image
sequence to sequence	a sequence	a sequence	a forecast at every step
encoder-decoder	a sequence	a sequence, after all of it	translation

Four shapes of recurrent model, and the cell is the same in all four. What changes is where you take the output from — the last state, every state, or a second network fed by the last state — and that choice is the architecture. We build the first and the third today; the fourth is Lecture 18.

The model is eleven lines — an `nn.RNN` and a linear head applied to `outputs[:, -1]`, the hidden state after the final step. The head is not optional: the state has 32 components and the target has one, and `tanh` outputs lie in $(-1, 1)$ while our scaled targets sometimes exceed that. `batch_first=True` because our batches are `[batch, time, series]`; without it the module silently reads the first dimension as time.

The training loop is Lecture 10's five lines, unchanged. The loss is Huber — quadratic near zero, linear far away — because we are optimising a proxy for MAE, and minimising the absolute error directly gives a gradient that never gets smaller as you approach the optimum. *The metric and the training loss are allowed to differ*: the metric is what the operator reads, the loss is what the optimiser can descend, and here is the case.

Shuffle, but shuffle the right thing

`DataLoader(train_set, shuffle=True)` shuffles the *windows*. It does not shuffle the values inside a window — the sequence within each row is untouched.

Consecutive windows share almost all their content: one starting on day t and one on day $t + 1$ differ in two values out of fifty-six. Gradient descent expects the instances in a batch to be independent draws. Without shuffling, a mini-batch is a run of near-identical, highly correlated windows: its gradient is effectively one example's, no less noisy than a single draw, and training stalls.

And do not let that shuffling cross the split. Windows are shuffled within the training period only. Shuffling and leakage are different operations, and it is easy to conflate them.

On identical shuffled folds the recurrent network scores 40,395, against 55,385 for copying last week and 44,925 for the linear model — 27.1% better than the baseline and 10% better than the linear model. Both beat the baseline comfortably. The recurrent model's lead over the linear one is the argument that it can use the *order* of the window rather than a weighted sum of it — but it is one comparison, with the recipe the ladder chose, and the ladder below shows how far one such choice can move a recurrent score.

Failure condition — and the assumption, stated rather than left implied

Those folds are shuffled, so *none of those numbers is a forecast*. Lecture 15 measured what that is worth. They are comparable with each other because the folds are identical, and with nothing else.

Every forward number in the rest of this lecture is on the honest ladder.

2 Deeper, wider, and further ahead

2.1 Loss, scale and what to watch

Squared error punishes a large miss quadratically, so on a series with occasional shocks a handful of days dominate the gradient. MSE is smooth everywhere and dominated by outliers; MAE is robust but its gradient is constant, so it converges slowly near the optimum; Huber,

$$L_{\delta}(e) = \begin{cases} \frac{1}{2}e^2 & |e| \leq \delta, \\ \delta(|e| - \frac{1}{2}\delta) & |e| > \delta, \end{cases}$$

has both properties for one knob. The choice is the same argument as choosing the metric in Lecture 15, with the same answer: it depends on what an error costs the operator.

Normalise by a constant computed on the training period only — the same rule as every other lecture, and easier to break here because the series looks like one object. Divide rather than standardise when the series is positive and multiplicative: ridership is counts, and a ratio is more natural than a difference. And if you differenced, un-difference the prediction before reporting it; the metric is in the units the operator uses.

Failure condition — a model scored in differenced units

A model evaluated on differences can look excellent and mean nothing: predicting a change of zero every day is nearly right in those units, and it is the naive forecast wearing a disguise. The units the metric was computed in decide what it is able to praise, and the report is usually silent about which they were.

While it trains, log the training and validation loss per epoch (the gap is Lecture 5's); the gradient norm — Lecture 11's probe, along time here, which on a plain RNN over a long window decays; and the forecast itself, plotted against the truth on a held-out stretch.

Failure condition — the failure a plot catches and a metric does not

A model that simply repeats the previous value scores well on almost any error metric and is worthless. On a plot it is unmistakable: the prediction is the truth, shifted right by one step.

2.2 Forecasting k steps ahead

Three strategies, all defensible, all to be evaluated the same way: forward in time, at the horizon you actually need. **Recursive** predicts one step, feeds it back and repeats, so errors compound and the model never trained on its own output. **Direct** fits a separate model per horizon: k models, none sharing what the others learned. **Sequence-to-sequence** is one model with k outputs at once: more to fit, and the outputs are not independent. A model that is excellent one step ahead and useless seven steps ahead is not a good model if the operator plans a week in advance.

Teacher forcing, and what it costs

Training a sequence-to-sequence model, you can feed the decoder the true previous token or its own previous prediction. The true token (**teacher forcing**) trains fast and stable, but the model has never seen its own mistakes; its own prediction trains slowly, early training is noise, and it matches deployment.

Failure condition — a model trained only on true prefixes

Trained on true history and evaluated on its own output, the model meets a distribution at prediction time it never saw in training — its own errors, compounding. This is **exposure bias**: the forecast drifts, the further ahead the worse, which is precisely the recursive strategy's weakness. And the training loss, computed on true prefixes throughout, never registers it.

2.3 Improvement 1 — deeper, which this run cannot rank

One argument, `num_layers=3`. On this run the simple RNN forecasts at 58,029 (train 37,388) and the three-layer one at 45,912 (train 31,359), both on the adam recipe: depth wins on test and on train.

Do not write that down. Each recipe is scored on July–December 2018, inside the training period, and the lower score wins. For the simple RNN this run scored sgd 41,878 and adam 40,910, so adam by 968 riders, 2.3%. The deck's original run scored sgd 41,878 and adam 42,879 — sgd by about 1,000, 2.3% the other way — and with sgd the simple RNN forecast at 49,073 against the deep RNN's 53,491, worse on train too (40,630 against 35,972). That run's reading was “not overfitting, a harder optimisation”. This run says the opposite with the same code.

The sgd score reproduced to the rider; the adam fits of both univariate RNNs did not, so the deep RNN moved as well on an unchanged recipe, while the linear and multivariate rows reproduced exactly. *A ranking that rests on one near-tied selection decision is not a finding.* Report it as unresolved, and settle it only with several seeds and a gap larger than their spread.

2.4 Improvement 2 — more series, which worked

We have been forecasting rail from rail alone while a bus series sits unused, and a `day_type` column we established is known in advance. Adding both — rail, bus, and the next day's type one-hot encoded, five columns — changes the architecture by one number, `input_size=5`.

Wait — that is a `shift(-1)`

Lecture 15 measured a random split that let a model train on rows surrounding what it was scored on. This one reads the future *deliberately*.

The test that separates them: *is the value available at prediction time?* The type of the day we are forecasting is on the public calendar and has been for a year. Its ridership is not.

This is the one question about leakage that no static check can answer for you. It is not a property of the code; it is a property of the world.

Model	Forecast MAE	vs copying last week
Copy last week	64,815	—
Simple RNN, rail only	58,029	10% better
Simple RNN, five inputs	40,046	38% better

Thirty-one per cent off this run’s rail-only error for one line of feature engineering and one changed constructor argument — eighteen against the original run’s rail-only row, which is the unstable one. The five-input row reproduced to the rider and beats both. Most of it is the day type: the network no longer has to infer that a holiday is coming from the shape of the previous eight weeks — which is impossible — because we told it.

2.5 Improvement 3 — further than one day

The operations team also wants a fortnight. Feed the forecast back — predict day $t + 1$, append it to the window as though it had happened, predict $t + 2$ — and errors accumulate. Or predict all fourteen at once: change the output layer to fourteen units and the target to a vector, one model and no accumulation.

A third way predicts fourteen at *every* step: at step 0 the forecasts for steps 1 to 14, at step 1 for steps 2 to 15, and so on, taking the head over all outputs rather than the last. That is not cheating: a recurrent cell at step t has seen steps 1 to t and nothing more — it is **causal**. The loss simply gains a term at every step, so many more gradients flow and they travel a shorter distance. After training, only the last step’s output is used.

Measured, the error grows with distance from 35,997 at one day to 46,586 at fourteen, and stays below the no-model baseline throughout.

3 LSTM and GRU, gate by gate

A plain cell multiplies the state at every step, so $\partial \mathbf{h}_t / \partial \mathbf{h}_{t-1} = \text{diag}(\phi') \mathbf{W}_h^\top$ and over T steps the gradient carries T copies of one matrix — Lecture 11’s argument with the extra force that it is one matrix rather than many independent ones.

What a gated cell does instead

An LSTM keeps a separate cell state updated *additively*,

$$\mathbf{c}_t = \mathbf{f}_t \odot \mathbf{c}_{t-1} + \mathbf{i}_t \odot \tilde{\mathbf{c}}_t.$$

When the forget gate is near one, $\partial \mathbf{c}_t / \partial \mathbf{c}_{t-1} \approx \mathbf{1}$ and the gradient passes through almost unchanged.

That is the entire trick: *a path along which the gradient is added to rather than multiplied.*

Gate	Form	What it decides
forget $\mathbf{f}_t$	sigmoid	how much of the old cell state to keep
input $\mathbf{i}_t$	sigmoid	how much of the candidate to write
candidate $\tilde{\mathbf{c}}_t$	tanh	what to write, if anything
output $\mathbf{o}_t$	sigmoid	how much of the cell state to expose as $\mathbf{h}_t$

Every gate is a function of the same two inputs ($\mathbf{x}_t, \mathbf{h}_{t-1}$) with its own weights, so a sigmoid in $[0, 1]$ is a *soft switch* and the network learns when to open it. Four gates means roughly four times the parameters of a plain cell.

A **GRU** has two gates. The update gate $\mathbf{z}_t$ merges the LSTM's forget and input gates into one — what is kept is what is not written; the reset gate $\mathbf{r}_t$ sets how much of the previous state to use when forming the candidate; and there is no separate cell state, $\mathbf{h}_t$ doing both jobs:

$$\mathbf{h}_t = (1 - \mathbf{z}_t) \odot \mathbf{h}_{t-1} + \mathbf{z}_t \odot \tilde{\mathbf{h}}_t.$$

Still additive, still a near-identity path when $\mathbf{z}_t$ is small — the property that mattered, kept, with about three quarters of the parameters. Which to use? On most problems they are within noise of one another: the GRU is smaller and trains faster, the LSTM has the longer track record. Measure, on your own protocol, rather than inheriting an opinion.

Model, all five inputs	Forecast MAE	Training MAE
Simple RNN	40,046	27,035
LSTM	37,542	24,540
GRU	34,639	20,356

The GRU wins, and the second column says why to be careful: it fits the training rows 25% better than the plain cell but the forecast only 14% better. That gap is overfitting, and it is what to attack next, before reaching for a bigger cell. Neither cell is a large model — a few thousand parameters. Across the whole ladder, every bar a forecast on the same 151 days, the error falls from 64,815 for copying last week to 34,639: 47% better than doing nothing.

4 The whole application

Forecast	What we reported	What it does
Copy last week	55,385	64,815
Linear, 56 lags	44,925	49,312
Simple RNN	40,395	58,029
GRU, five inputs	—	34,639

The two middle rows are the same models measured twice; the simple RNN's right-hand number is the one that moves between runs. The linear row's left number is Lecture 15's scikit-learn fit; its right number is this lecture's, on the same 151 days as every other right-hand number. The bottom row is the model we would deploy, and it was never measured the wrong way. *The final number is better than the one first reported — but we reached it by first admitting that the first one was not a forecast.*

Is 34,639 good? It is 47% better than a system a transit authority could run with no data scientist at all, and its typical miss is around 10.4% of the day it forecasts, against 16.5% for copying last week *on those same 151 days*. A naive percentage measured over any other set of days is not the comparison to make here. (As a share of a mean day on those same 151 days — 583,030 riders — the miss is 5.9%, which is a third quantity again.)

What a defensible report says

“Forecasting the five months after the training cut, and never seeing them, the model’s mean absolute error is 34,639 boardings a day against 64,815 for copying the same day of the previous week. Both figures are on identical rows. The gain is concentrated on holidays.” The protocol, the baseline, the rows, and where the gain comes from — four things the number alone does not carry.

4.1 Red-team it: what happens when the world changes

Fit on everything to the end of 2019 and forecast the four months from March 2020.

	MAE
The linear model, trained to the end of 2019	84,380
Copying last week, on the same 122 days	36,282
The actual mean level over that period	141,404

Failure condition — a regime the training window never contained

Scored on the four months from March 2020, the model is 2.3 times worse than doing nothing and its error is six tenths of the entire level of the series. Nothing in training could have shown this: the failure is that the future stopped resembling the window the model was fitted on.

To be fair to the model: March 2020 is not a leak and not its fault. No forecaster could have produced it from eight weeks of history, and copying last week wins only because it adapts within a week of the collapse. The lesson is not that the model is bad — it is that *an honest protocol told you the model had failed, on the day it failed*. A shuffled split, with those same days spread across its training folds, could not have. Which is why a deployed forecaster is monitored against its own naive baseline, continuously: if the margin disappears, something has changed in the world.

Temporal leakage checklist

1. Is the split by *date*, not by row?
2. Does every feature exist at the moment the forecast is made? Name the moment.
3. Were any statistics — means, scalars, encodings — computed over the whole series before the cut?
4. Do the training windows overlap the scored windows? Purge them.
5. Is the baseline computed on identical rows to the model?
6. Does the reported score come from more than one origin?

The third is the oldest leak in the course — a statistic fitted before the split — and the fourth is the one only a series of overlapping windows has.

Reading a recurrent result honestly

At what horizon?	one step ahead is a different problem from seven
Against which baseline?	seasonal-naive, not the mean — Lecture 15
On which protocol?	rolling origin, and how many folds
With what spread?	one origin is one sample
In which units?	differenced or not; scaled or not

A recurrent network that beats the naive forecast by less than the fold-to-fold spread has not beaten it — and a ranking that flips between runs, as the depth comparison did, was never a ranking.

A recipe for a forecasting model

1. Plot the series. Seasonality and trend are usually visible, and cost nothing to find.
2. Compute the autocorrelation, and the naive and seasonal-naive baselines. These are the numbers to beat.
3. Fix the protocol before fitting anything: horizon, rolling origin, number of folds, metric.
4. Fit a linear model on lags. It is fast, and it is often close.
5. Only then a recurrent model, scored on the same protocol.
6. Report the ladder with its spread, not the best row alone.

Steps 2 and 3 come before step 4 because after you have a number, you will not choose the protocol that makes it look worse.

5 The notebook

The notebook builds the simple recurrent model line by line, then runs the whole ladder of six models under one protocol: each recipe chosen on July to December 2018, inside the training period, and every model scored on the same 151 forward days. It also scores the simple RNN on shuffled folds, forecasts fourteen days ahead, and fits the linear model through March 2020.

What to change

Change the seed in the setup cell and re-run the ladder. Watch the simple and the deep RNN: if their order flips, you have reproduced the lecture's point about a ranking that rests on one near-tied choice.

6 Exercises

The five set at the end of this lecture's deck, in the style of the written paper. The solutions are on Lecture 17's deck and in the exercise book.

1. A recurrent network is trained on 56-day windows drawn from one series. Explain why shuffling the *windows* is legitimate while shuffling the *series* is not. [4]
2. A three-layer RNN beats a one-layer RNN on test and on train. A re-run of the same code reverses the order on both. State what the ranking rested on, and what would settle it. [4]
3. Gates were introduced after a plain recurrent cell failed over long windows. State the mechanism, in terms of what happens to the same matrix over 56 steps. [4]
4. A colleague builds the multivariate forecaster with `day_type.shift(-1)` as a feature and objects that it reads the future. State whether it is a leak, and the one question that settles it. [5]
5. Report the training MAE beside the test MAE. Name the two failures the pair distinguishes, and say why the test column alone cannot. [4]

Summary

A recurrent cell is Lecture 9's layer with one extra term, and that term decouples reach from parameter count — weight sharing along time instead of across an image. The state is a fixed-width lossy summary, which is a budget: if the task needs more bits than it holds, no amount of training supplies them. Unrolling makes training ordinary backpropagation on a 56-layer network with *shared* weights, and that sharing is the problem, because the gradient then carries a power of one matrix rather than a product of independent ones.

Three improvements, measured forward in time. Depth is the one that cannot be ranked: on this run the three-layer network beat the simple one on test and on train, on the deck's original run it lost on both, and the difference is one near-tied choice of training recipe that went the other way. A ranking that rests on such a choice is not a finding. Adding the bus series and the next day's type — a `shift(-1)` that is legitimate because the calendar is public — paid more than depth did, and reproduced. Gates buy an additive path along which the gradient is added to rather than multiplied, and the GRU on five inputs reached 34,639.

The ladder ends 47% better than copying last week, and the honest ladder is not the one we first reported: the two middle rows are the same models measured twice. Then the red team — a model trained to the end of 2019 is 2.3 times worse than doing nothing on the months from March 2020. An honest protocol reports that failure on the day it happens; a forecaster in deployment is monitored against its own naive baseline for exactly that reason.