

Convolutional networks

Flowers102

LECTURE 12

GÉRON CH. 12

Applications of Machine Learning

BSc Mathematics of Artificial Intelligence · Third year

Where we are

Lecture 11 took a deep fully connected stack on flattened 28×28 greyscale images that would not train, and repaired it:

- Two probes located the fault: the forward signal looked healthy, the *backward* one shrank by a constant factor per layer
- One variance recursion explained it, and Glorot's and He's initialisations are the two ways of setting that factor to one
- Six repairs — initialisation, activation, batch normalisation, clipping, optimiser, schedule — added one at a time, and an ablation saying which of them paid

WHAT BREAKS WHEN THE IMAGE GETS BIGGER

Flattening a 224×224 colour image gives 150,528 inputs. A single dense layer of 1,000 units on that is **150 million parameters** — for one layer, on one image size. Every one of them has to be learned from data, and none of them knows that two adjacent pixels are related.

Today's method fixes both problems at once, and the derivation says exactly how much it saves.

Today

1. The task, and why a flattened image is the wrong representation
2. **The mathematics** — what weight sharing saves, equivariance, and how pooling buys invariance
3. **The method** — convolution, pooling, an architecture, and training it
4. What it learned, what it scores, and what goes wrong quietly

[Open Lecture 12 in Colab](#)

The stakeholder

A horticultural wholesaler receives mixed pallets of cut stems and has to sort them by species before they reach the auction floor.

- One overhead camera per lane, one photograph per stem
- A human inspector currently does this and is the bottleneck
- They want the machine to name the species from the photograph

Same shape of problem as the garment sorter in Lecture 9. Everything that makes it harder is in the pixels.

What they actually asked for

IN THEIR WORDS

“Tell us which of our species it is. If you cannot tell, say so and we will look at it ourselves — but do not tell us the wrong one.”

Two requirements hiding in one sentence: a classifier over a fixed list of species, and a usable notion of *confidence*.

We build the first today. The second waits until the numbers are good enough to be worth calibrating.

FAILURE CONDITION — A CLASSIFIER WITH NO WAY TO ABSTAIN

The stakeholder asked to be told when the model cannot tell. A softmax always returns a class, and its largest probability is not a confidence you have calibrated. Deliver it as one and you have answered a different requirement while appearing to meet this one.

The dataset

`torchvision.datasets.Flowers102` — 8,189 photographs of 102 flowering species common in the United Kingdom.

```
from torchvision.datasets import Flowers102

train = Flowers102("data", split="train", download=True)
val    = Flowers102("data", split="val",    download=True)
test   = Flowers102("data", split="test",   download=True)
```

About 345 MB. It is one of the datasets the textbook names for this chapter, and the split used here is the one shipped with it — so every result in the literature is on the same 6,149 test images.

What the camera sees

12 of the 102 species, one image each, resized to 128×128


daffodil



mallow



azalea



sunflower




cautleya spicata



hibiscus



monkshood




bougainvillea




globe thistle




windflower



blanket flower



giant white aru...



Different scales, different backgrounds, different lighting — and several of these species are hard to tell apart by eye.

What is hard here, precisely

- **Position** — the flower is not centred, and not at a fixed scale
- **Background** — leaves, soil, other flowers, a gardener's hand
- **Between-class similarity** — several of the 102 species are separated by petal count and little else
- **Within-class variation** — the same species at two stages of opening

Only the first of those has a structural answer. It is the answer this chapter is about.

The split we are given

Split	Images	Per class	Used for
train	1,020	10, exactly	fitting
val	1,020	10, exactly	choosing
X test	X 6,149	20–238	reporting, once

X There are six times as many test images as training images.

Why the training split is so small

Because somebody had to look at every photograph and name the species.

THE STANDING CONSTRAINT OF APPLIED VISION

Images are cheap and labels are not. Every application in this part of the course has more pixels than it can use and fewer labels than it wants.

Design for that. A method that needs thousands of labelled examples per class is not a method you can deploy here.

The obvious thing to try first

You have a multilayer perceptron that works. Resize every photograph to 128×128 , flatten it, and feed it in.

	Count
Inputs, $128 \times 128 \times 3$	49,152
Outputs, 32 values at every pixel position	524,288
X Weights in one fully connected layer between them X 25,769,803,776	

X That is 96 GB of float32 for one layer. The full argument is this lecture's mathematics block; here it is just a reason to do something else.

What weight sharing saves

Input $3 \times H \times W$, output $32 \times H \times W$. A dense layer connects every output to every input: $n_{\text{in}}n_{\text{out}} = 96 H^2 W^2$ weights. A convolution connects each output to a $k \times k$ window of every channel, **with the same weights at every position**: $C_{\text{out}}C_{\text{in}}k^2 = 32 \times 3 \times 7 \times 7$, which is 4,704 weights.


$$\frac{n_{\text{in}}n_{\text{out}}}{C_{\text{out}}C_{\text{in}}k^2} = \frac{H^2W^2}{k^2}$$

5,478,275

times fewer weights at $H = W = 128$, for the same output shape.

With the 3×3 filters used everywhere else in the network the factor is larger again, by $(7/3)^2$: the dense count never contained k .

The same output shape, two ways



A logarithmic axis, and the two bars still barely fit on the same figure.

Equivariance, precisely

Sharing the weights assumes that a pattern means the same thing wherever it appears. Let $T_{\mathbf{v}}$ be the operator that shifts an image by the vector $\mathbf{v}$:

$$(T_{\mathbf{v}}x)(\mathbf{p}) = x(\mathbf{p} - \mathbf{v})$$

A map f is **equivariant** to translation when

$$f(T_{\mathbf{v}} x) = T_{\mathbf{v}} f(x) \quad \text{for every } \mathbf{v}$$

Shift then compute, or compute then shift — the same answer. A convolution has this property because its weights do not depend on the position; it holds exactly away from the border, where zero padding invents input that was not there.

How pooling manufactures invariance

A map g is **invariant** to translation when

$$g(T_{\mathbf{v}} x) = g(x) \quad \text{for every } \mathbf{v}$$

The output does not move at all, so it no longer says *where*. Classification wants that; equivariance alone does not give it.

A 2×2 max pool is invariant to any shift that keeps the maximum inside its own window — so half the one-pixel shifts, in each direction.

- Stack four pooling stages and the tolerated shift grows to ± 8 pixels in the last feature map
- A global pool over the whole map is invariant to *every* shift — there is no position left to be sensitive to

The invariance is approximate and it is built, one factor of two at a time. Nothing in the architecture declares it.

The idea, in one sentence

A feature worth detecting in the top left corner is worth detecting in the bottom right, so **use the same weights everywhere**.

- That is the whole of it. Everything else is bookkeeping
- It buys two things at once: far fewer parameters, and a network that does not have to relearn a petal edge at every position

Both halves were stated in the mathematics block. Here we use them.

A convolution, as arithmetic

For an input x with C channels and a kernel w of size $k \times k$, the output at position (i, j) of filter f is

$$z_{ij}^{(f)} = b^{(f)} + \sum_{c=1}^C \sum_{u=0}^{k-1} \sum_{v=0}^{k-1} w_{c,u,v}^{(f)} x_{c,i+u,j+v}$$

A dot product between the kernel and a window of the input, repeated at every position. The sum over c is why a filter reads *all* the input channels at once.

`nn.Conv2d` , and the defaults you did not ask for

```
1  nn.Conv2d(  
2      in_channels=3,          # must match the previous layer's output  
3      out_channels=32,        # how many filters -> how many output maps  
4      kernel_size=7,  
5      stride=1,               # the default  
6      padding=3,              # NOT the default; the default is 0  
7      bias=False,             # a batch-norm follows, and it has its own shift  
8  )
```

The default you did not ask for: **padding defaults to zero**, which silently shrinks every feature map by $k - 1$ and will eventually give you a shape error twelve layers later.

Pooling

After a couple of convolutions, halve the height and width by keeping only the largest value in each 2×2 block.

- Three quarters of the activations disappear, and with them three quarters of the memory and the arithmetic in every later layer
- It has **no parameters** — nothing about it is learned
- Each later unit therefore sees a larger piece of the original image

There is a fourth consequence — invariance — and it is the one the mathematics block ended on.

Max pooling, on numbers

4 × 4 feature map

5	3	8	1
2	9	4	0
1	6	7	3
4	2	5	8

after 2 × 2 max pool

9	8
6	8

Sixteen numbers in, four out, and the four are the largest of each block.

Why every convolution is followed by a batch norm

Because ten layers of untreated convolutions on 1,020 images do not train. This is Lecture 11's $\rho \neq 1$ in a new architecture — and here, unlike there, initialisation alone does not reach it. Lecture 11 measured batch normalisation making a *repaired* network worse by 3.2 points; the size of an effect depends on what else is in the system.

- Ten layers of untreated convolutions on 1,020 images do not train: the loss sits at $\log 102 \approx 4.62$ and the network predicts one class
- `BatchNorm2d` normalises across the batch *and* both spatial axes, so it holds one mean and one variance per channel

```
nn.Conv2d(c_in, c_out, 3, padding=1, bias=False)
nn.BatchNorm2d(c_out)      # 2 * c_out parameters, not 2 * c_out * H * W
nn.ReLU()
```

The architecture

Stage	Layers	Output
1	conv 7×7 ×32, conv 3×3 ×32, pool	32 × 64 × 64
2	conv ×64, conv ×64, pool	64 × 32 × 32
3	conv ×128, conv ×128, pool	128 × 16 × 16
4	conv ×256, pool	256 × 8 × 8
✓ head	✓ flatten, dense 256, dropout, dense 102	✓ 102

Channels double as the map halves — the textbook’s own rule of thumb: what pooling takes away in resolution, the extra filters give back in channels.

The architecture, in code

```
1  def conv_block(c_in, c_out, k=3):
2      return [nn.Conv2d(c_in, c_out, k, padding=k // 2, bias=False),
3              nn.BatchNorm2d(c_out), nn.ReLU()]
4
5  net = nn.Sequential(
6      *conv_block(3, 32, k=7), *conv_block(32, 32), nn.MaxPool2d(2),
7      *conv_block(32, 64),      *conv_block(64, 64), nn.MaxPool2d(2),
8      *conv_block(64, 128), *conv_block(128, 128), nn.MaxPool2d(2),
9      *conv_block(128, 256),      nn.MaxPool2d(2),
10     nn.Flatten(),
11     nn.Linear(256 * 8 * 8, 256), nn.ReLU(),
12     nn.Dropout(0.5),
13     nn.Linear(256, 102),
14 )
```

Thirty modules. Nothing in it that was not on a slide in this lecture or the previous one.

Check the shapes before you check anything else

```
x = torch.zeros(2, 3, 128, 128)
for m in net:
    x = m(x)
    print(f"{type(m).__name__:12s} {tuple(x.shape)}")

assert x.shape == (2, 102), f"the head is wrong: {x.shape}"
```

What is the shape here? Asking costs four lines. A shape error inside a training loop surfaces as a size-mismatch traceback two hundred lines deep; here it surfaces as a printed line.

Count the parameters, and look at where they are

Part	Parameters	Share
Every convolution and batch-norm layer together	586,720	12.2%
X One dense layer, 16,384 → 256	X 4,194,560	X 87.3%
Output layer, 256 → 102	26,214	0.5%
✓ Total	✓ 4,807,494	100%

X Nine parameters in ten are in the layer that is not convolutional.

The training loop — unchanged

```
opt    = torch.optim.Adam(net.parameters(), lr=3e-4)
lossf  = nn.CrossEntropyLoss()

for epoch in range(30):
    net.train()
    for xb, yb in train_loader:
        opt.zero_grad()
        loss = lossf(net(xb), yb)
        loss.backward()
        opt.step()
```

Exactly the five lines from Lecture 10. Convolution changed the model; it did not change the loop, because autograd does not care what the modules are.

One hyperparameter that is not a free choice

Learning rate	Seed 42	Seed 43	Mean
10^{-3} , Adam's default	0.98%	1.86%	1.42%
3×10^{-4}	6.86%	4.61%	5.74%

Validation accuracy after 12 epochs, on this architecture and this data. The default does not diverge — it plateaus at or near chance, which on a split with ten images of every species is 0.98%, and that is the failure mode that does not announce itself.


4.31 points, from one keyword argument nobody wrote down — about twice the 2.25 points between the two seeds of the good rate.

Run it

Setting	Value
Epochs	30
Batch size	32
Steps per epoch	32
Total optimiser steps	960
X Wall clock, seed 42 X 62 s	

Measured on an Apple MPS backend, no CUDA, and including a validation pass after every epoch. A free Colab T4 is in the same range for a model this size; the point of the number is the comparison we make next lecture, not the hardware.

The cost, on the x-axis



The x-axis is seconds, not epochs. That is deliberate: from the next lecture on, time is one of the things being compared.

The test set. Once.

```
net.eval()
with torch.no_grad():
    correct = sum((net(xb).argmax(1) == yb).sum().item()
                  for xb, yb in test_loader)
print(f"test accuracy {correct / len(test_ds):.4f}")
```

11.71%

One run, seed 42. Two seeds of a shorter run of the same network land 2.25 points apart, so one run of a network this size on 1,020 images is an anecdote.

Why one run is not a result

Seed	Validation accuracy, 12 epochs
42	6.86%
43	4.61%

X 2.25 points between two runs, and nothing changed but the seed.

Any comparison you make from here that is smaller than that is not a comparison. Two seeds only show that the spread exists; a mean you would report needs more runs than a free CPU affords.

One seed reported as a result


FAILURE CONDITION

2.25 points separate two runs with nothing changed but the seed. A single run is one draw from that spread, and a comparison of two single runs is two draws — which is why what gets reported is the spread as well as the mean.


Thirty-two filters, before and after

All 32 filters of the first layer, $7 \times 7 \times 3$, each rescaled to its own range

at initialisation



after 30 epochs



Look at the two blocks for ten seconds, and say what changed.

It moved a little. It did not become anything

```
import torch.nn.functional as F

# init_filters: a copy of net[0].weight, taken before the first step
before, after = init_filters.flatten(), net[0].weight.flatten()
print(f"cosine similarity {F.cosine_similarity(before, after, dim=0):.4f}")
print(f"mean |change| / mean |weight| "
      f"{(after - before).abs().mean() / before.abs().mean():.1%}")
```

	Value
✓ cosine similarity, all 4,704 weights before against after	0.9880 ✓
mean absolute change, as a share of the mean weight	14.1%
weights that changed sign	160

X The weights moved. No structure appeared — the right-hand block is still noise, and 62 seconds of training bought it.


Two reasons, and both are worth knowing

1. **The gradient that reaches it is small.** Ten layers of chain rule separate the loss from these weights, and 1,020 images give 960 optimiser steps in total
2. **Its scale does not matter.** A batch norm follows immediately, and normalising the output makes the overall size of these weights irrelevant to the function — so there is no gradient pressure on it at all

The textbook picture of a first layer — colour blobs, oriented edges — is real, and you will see it next lecture. It comes from a corpus a thousand times this size.

What a filter does to one photograph

Feature maps: the first layer keeps the picture, the fourth keeps a summary



Early maps still look like the photograph. Four layers down they do not, and that is the point.

Reading a feature map

- Bright means *this filter fired strongly here* — nothing more mystical than that
- The spatial layout survives: a bright patch is where in the image the pattern was found
- After four pooling stages the map is 8×8 , so one cell answers for a 256-pixel square of the original

That last line is a loss of information we chose. It is the right choice for naming a species, and the wrong one for any task that must say *where*.

The scoreboard

	Test accuracy
Uniform guess	0.98%
Commonest species	3.87%
✓ Yours, today	11.71% ✓
Operator's requirement	90%

Far above both baselines, and far below the requirement. Both distances matter.

What this number is and is not

- **Is:** one run, seed 42, of one architecture on one split, measured on 6,149 held-out images
- **Is not:** a claim about convolutional networks. It is a claim about this one, trained on 1,020 images
- Two seeds of a shorter run differed by 2.25 points. Any comparison smaller than that is not a comparison

Quote the spread whenever you quote the number. It is the difference between a measurement and an anecdote.

Five ways this build goes wrong quietly

Mistake	Symptom
<code>padding=0</code> left at the default	maps shrink; a shape error many layers later
Forgetting <code>net.eval()</code>	the score wobbles between calls — Lecture 10
Normalising with statistics of all splits	0.20 points, inside the seed noise
<code>lr</code> left at Adam's default	4.31 points, and no error
Pooling before any convolution	throws away the pixels before anything reads them

The notebook for this lecture

[Open Lecture 12 in Colab](#)

- **Runs on free CPU.** It trains 30 epochs, and every accuracy on these slides is one it printed
- The dense-against-convolutional parameter count, computed rather than quoted — the arithmetic you can do on paper first
- The network built and shape-checked layer by layer, the learning-rate and seed comparisons, one test score, and the first layer's filters and feature maps

WHAT TO CHANGE

Set `padding=0` on every convolution and re-run the shape check. Every map shrinks and the first linear layer no longer matches.

What we did

1. Counted the weights of a dense layer and a convolutional layer producing the same output, and found the image size absent from the convolutional count
2. Defined translation *equivariance*, the property weight sharing buys, and saw that pooling builds *invariance* on top of it
3. Built a convolutional network — convolution, batch norm, pooling, a dense head — checked its shapes, and found nine parameters in ten in the one dense layer
4. Trained it with a learning rate that is not a free choice, scored the test set once, and quoted the seed spread beside the number
5. Looked inside: the first-layer filters barely moved, and the feature maps lose the photograph as they go down the stack

NOT PRESENTED — FOR AFTERWARDS

Exercises

Lecture 12

five, in the style of the written paper

Exercises — Lecture 12

- 1** A dense layer mapping a $3 \times 128 \times 128$ input to a $32 \times 128 \times 128$ output has $n_{\text{in}}n_{\text{out}}$ weights. Give the convolutional count for a 7×7 kernel, and say which symbols are absent from it. [5]
- 2** State the difference between equivariance and invariance, say which of the two pooling builds, and why a species classifier wants it. [4]
- 3** Convolution is equivariant to translation on the interior of an image but not at the border. Give the reason. [3]

Solutions on Lecture 13's deck.

Exercises — Lecture 12 (continued)

- 4** Trained with Adam at its default learning rate of 10^{-3} , the lecture's network neither diverges nor raises an error. State what it does instead, and how the lecture showed it. [4]
- 5** Two runs of the same network, differing only in the seed, reach 6.86% and 4.61% validation accuracy. A change to the pipeline then gains 1.5 points in a single run. What can you conclude, and what should be reported? [4]

Solutions on Lecture 13's deck.

THE FIVE SET AT THE END OF LECTURE 11

Solutions

Lecture 11

not part of this lecture

Solutions — Lecture 11

1. One layer multiplies the standard deviation of the backward signal by...

✓ They multiply it by $\rho^{L-1} - L - 1$ steps between the first gradient and the last — so anything but 1 vanishes or explodes geometrically.

- A constant applied L times is exponential in L
- There is no regime where a repeated constant other than 1 is safe

2. Preserving the forward signal asks for $\text{Var}(w) = 1/n_{\text{in}}$ and preserving the gradient...

✓ They agree only when $n_{\text{in}} = n_{\text{out}}$; Glorot takes the harmonic mean, $2/(n_{\text{in}} + n_{\text{out}})$.

- Any layer that changes width cannot satisfy both
- The compromise accepts a small error in each direction rather than a large one in either

Solutions — Lecture 11 (2)

3. He initialisation doubles the weight variance for ReLU. Derive the factor of two in one line.

✓ **ReLU zeroes half of a symmetric zero-mean input, so**

$$\mathbb{E}[\text{ReLU}(z)^2] = \frac{1}{2}\mathbb{E}[z^2].$$

- Exactly half the variance is discarded at each layer
- Doubling $\text{Var}(w)$ puts it back

4. A gradient profile that is a straight line on a log axis tells you something a curved one does not. State...

✓ **The attenuation is the same factor at every layer, so it is the initialisation rather than one bad layer.**

- A constant ratio per layer is a straight line in log space
- A single broken layer would show a step, not a slope

Solutions — Lecture 11 (3)

5. Gradient clipping is applied to a network whose gradients are fifteen orders of magnitude too small. Predict...

✓ No effect — clipping bounds gradients that are too large.

- The threshold is never reached, so nothing is rescaled
- Reaching for it means the failure mode was never measured

LECTURE 13 · GÉRON CH. 12

Transfer learning

Why features learned on a large labelled corpus transfer to these flowers
Freeze the backbone and train a head; then let the last block move, at its own learning rate
Augmentation — and why it belongs to the training loader only

Run the Lecture 12 notebook first